

The Constant in Front of $\sqrt{n}$: Measured Pollard Rho, ECC2K-130 Arithmetic, and Index Calculus over Prime Fields

Anonymous

Abstract

This library measures discrete-logarithm algorithms in one unit $S = \text{ops}/\sqrt{n}$, against two boundaries: a counting floor and a matched Pollard rho. On elliptic curves over prime fields and over $\mathbb{F}_p$, the method sits hundreds of times above rho at every size that fits; the constant in front of $\sqrt{n}$ is the quantity that moved, and it did not fall below one. The same unit prices a multiplicative-group linear sieve that does cross rho near 56 bits, a Koblitz GPU/FPGA iteration for ECC2K-130 whose pipe floor rules out 28 billion steps per second on one GPU, and a genus-3/4 Jacobian positive control that behaves as the exponent law says.

Contents

1	Introduction	3
1.1	What a number is allowed to mean	3
1.2	Two repositories	4
1.3	What is not claimed	4
2	Method: boundary, unit, table, class	4
2.1	Two boundaries	4
2.2	One table, one unit	4
2.3	Four classes	5
2.4	Worked case: residual walks	5
2.5	Worked case: the ECC2K-130 pipe	6
2.6	What does not count	6
3	Pollard rho	6
3.1	The walk	6
3.2	Negation map and fruitless cycles	6
3.3	Measured S	7
3.4	Fitted exponents	7
3.5	GPU kernel, emulator-verified	8
3.6	Fleet protocol	9
4	Square-root neighbours, Cheon, and cross-validation	9
4.1	Two grumpy giants and a baby	9
4.2	Free precomputation	9
4.3	GLV folding	10
4.4	Cheon	10
4.5	Cross-validation against the companion	10

5	ECC2K-130 arithmetic	11
5.1	Why a type-II optimal normal basis	11
5.2	Three multiplier realisations	12
5.3	Two iteration functions	12
5.4	Golden model and CPU walker	13
5.5	Small-curve controls, including ECC2K-95	13
6	ECC2K-130 on GPU	13
6.1	The ledger: 0.85 to 14.64 B/s	13
6.2	The floor that rules out 28 B/s	14
6.3	Table walk, declined at 16.56 against 17.0	14
6.4	Per-GPU survey	15
6.5	This library’s imported client	15
7	ECC2K-130 on FPGA	15
7.1	Digit-serial core	16
7.2	Companion engine on F2	16
8	ECC2K-130 campaign and structural negatives	17
8.1	Cutoff 32, $2^{28.41}$ iterations per point	17
8.2	Certification and the 124 M-point snapshot	18
8.3	Structural negatives	18
9	Index calculus over prime fields	19
9.1	Linear sieve in $(\mathbb{Z}/p\mathbb{Z})^*$	19
9.2	Elliptic curves: the negative, phase by phase	20
9.3	$E(\mathbb{F}_{p^3})$ residual walks	20
9.4	Joux–Vitse pair-only	21
9.5	Genus 3/4 Jacobian, the positive control	21
10	Related work	21
11	Conclusion	22
A	Reproduction commands	23
A.1	This library	23
A.2	Sibling study library	23
A.3	This paper	24

1 Introduction

The discrete logarithm in a group of prime order n has a generic cost of $\Theta(\sqrt{n})$ group operations. Pollard’s rho method makes the leading constant about $\sqrt{\pi/2} \approx 1.25$ without automorphisms and about $\sqrt{\pi/4} \approx 0.89$ with the negation map [5, 19, 22]. Every subsequent algorithm that claims to be faster is a claim about that constant, or about the exponent. This paper measures both.

The work is a study library, not an attack on cryptographic sizes. Field arithmetic is 64-bit Montgomery arithmetic; every problem the library solves lives in a group of order below 2^{64} . That bound is deliberate. The purpose is to put every variant on one axis, with every phase inside the cost, and to read off whether the constant in front of $\sqrt{n}$ actually moved.

Three families are priced.

Pollard rho, and the square-root neighbours. Teske’s r -adding walk, the negation map with Bernstein–Lange–Schwabe cycle handling, one inversion per 256 additions, parameters chosen from n . Baby-step giant-step, two grumpy giants and a baby, Pollard kangaroo, Bernstein–Lange free precomputation, GLV folding, Cheon’s auxiliary-input attack, a CUDA kernel verified through a host emulator, and a fleet protocol. The unit

$$S = \frac{\text{group operations}}{\sqrt{n}}$$

makes rho a flat line. Everything else is a row.

ECC2K-130. The Certicom Koblitz challenge $y^2 + xy = x^3 + 1$ over $\mathbb{F}_{2^{131}}$, group order $4 \cdot r$ with r a 129-bit prime, expected $2^{60.9}$ iterations [1, 6]. Optimal-normal-basis arithmetic, three multiplier realisations, a golden C model, a bitsliced client, a packed CUDA walk measured from 0.85 to 14.64 billion iterations per second, a measured pipe floor that rules out 28 B/s on one GPU, a table walk declined at 16.56 against a pre-declared 17.0, a per-GPU survey, a digit-serial FPGA core in this library (362 cycles/step, 7,941 LUT4) and a companion engine measured on an AWS F2 at 8.25 G steps/s, and a live campaign at cutoff 32 with $2^{28.41}$ iterations per point, about 124 million points, and no collision.

Index calculus over prime fields. A Coppersmith–Odlyzko–Schroeppel linear sieve in $(\mathbb{Z}/p\mathbb{Z})^*$ with Lanczos and Hensel lifting, whose crossover with rho is near 56 bits and which we re-ran through 60 and 62 bits. Then the elliptic negative: Semaev summation polynomials on $E(\mathbb{F}_p)$ and residual walks on $E(\mathbb{F}_{p^3})$, priced phase by phase with fitted exponents, a Joux–Vitse pair-only row, oracle rounds labelled as stage diagnostics, and a genus-3/4 Jacobian positive control whose ratio to rho falls as the exponent law predicts.

1.1 What a number is allowed to mean

A ratio to a boundary is the result; a constant that moved while the ratio stayed flat is engineering; a headline count that fell while total S rose is relabelling; a number that changed because the accounting did is a correction. Section 2 states the rule and works two cases. The residual-walk thread is the cautionary one: its relation phase crosses rho near 2^{96} in isolation, while the whole method never does, because the linear algebra it had left out grows as $n^{0.68}$ against relations at $n^{1/3}$.

Wall-clock time is a practicality note. Operation counts survive hardware. A run whose answer was not checked against a planted secret is not a result. Companion CUDA rho over 256-bit primes, the ECC2K-95 kernel, and kangaroo code whose READMEs record that they have run only through CPU emulation are kept out of headline claims; they appear as cross-checks of the walk, not as device throughputs.

1.2 Two repositories

The measurements span two checkouts. This library (`cryptanalysis`) is the 64-bit C11 reference: generic cyclic groups, the solvers of Section 3, the ECC2K-130 golden model and digit-serial FPGA core, the imported packed GPU client held to that model, and the $(\mathbb{Z}/p\mathbb{Z})^*$ sieve. The sibling study library (`crypto`), named from `docs/ALGORITHMS.md`, is where the ECC2K-130 campaign, the GPU ledger, the F2 FPGA engine, the residual-walk and hyperelliptic notes, and the index-calculus scoreboard live. Paths in footnotes without a prefix are in this library; paths beginning `crypto:` are in the sibling. The paper covers both and says so here, once.

1.3 What is not claimed

No elliptic-curve index-calculus variant measured end to end sits below matched rho. No one-addition-per-step walk on ECC2K-130 reaches 28 B/s on one GPU. No Modal SKU in the survey beats the RTX PRO 6000 at the campaign recipe. The 256-bit CUDA rho of `crypto:gpu/ecc/` and the ECC2K-95 kernel of `crypto:gpu/ecc2k/` are emulator-verified walks, not device-throughput results. Extrapolations are marked as such and carry the exponents they rest on.

2 Method: boundary, unit, table, class

A thread that claims to improve an attack must express its progress as a ratio to a boundary, reported in a single table whose rows are variants and whose columns are one unit. A thread that cannot state its boundary has not started; a thread whose ratio to the boundary is flat has not progressed, however much its constants moved. The rule is the reporting convention of the sibling study library.¹

2.1 Two boundaries

A *boundary* is a quantity the method under study cannot cross, derived rather than measured. A thread normally has one of each kind.

- A *floor*: a lower bound from a counting or generic-group argument. For residual walks on $E(\mathbb{F}_{p^3})$ a generic algorithm’s expected relation yield from P residuals is at most $\gamma P^2/2n$, so the count factor obeys $\kappa_{\text{total}} \geq \sqrt{2(B+1)/\gamma}$. That floor moves only with the factor-base size B and the automorphism order γ ; it cannot be tuned away.²
- A *reference*: the cost of the best algorithm that already solves the same problem, measured in the same unit on the same instances. For the discrete logarithm that is Pollard rho, run on the same group, with the same operation accounting.

Both are written down *before* optimising. A boundary chosen after the fact is not a boundary.

2.2 One table, one unit

The unit used throughout is

$$S = \frac{\text{total operations}}{\sqrt{n}}.$$

Rho is then a flat $S \approx 1.3$ at every size, and “is this better than rho” is one column. Every cost the method incurs belongs in S : precomputation, table setup, relation verification, linear

¹`crypto:AGENTS.md`.

²`crypto:research/notes/index-calculus/RESEARCH_RESIDUAL_WALKS.md`, §9.6.

algebra, and any work an oracle does per call. Foreign units are converted with a measured factor (for instance 63 field multiplications per curve addition on $\mathbb{F}_{p^3}$) and the factor is recorded.

End-to-end speed is the measure of speed. S is defined over the whole method, cold, from setup to the recovered logarithm. A number that prices only one phase — the relation search, the decomposition oracle, a single solver call, the linear algebra — is a *stage diagnostic*, never a speed. “Faster than rho” means the method’s S column, with every phase inside it, sits below rho’s, robustly, at growing n . The only admissible speed ratio is

$$\text{speedup} = \frac{\text{baseline total operations}}{\text{candidate total operations}}.$$

A ratio taken over any smaller slice of the pipeline is labelled a stage diagnostic and is not reported as a speedup. A row without a verified answer is not a result.

2.3 Four classes

Every change is classified by what moved, not by how it felt to make.

class	what moved	what it means
advance	ratio to the floor fell	the method does something a generic algorithm cannot
engineering	S fell, ratio to the floor flat	legitimate, bounded, and not a finding
relabelling	headline count fell, S rose	work moved somewhere the headline does not look
accounting	numbers changed, algorithm did not	a correction; say so and do not claim the gain

All four are worth committing. Only the first is a result.

2.4 Worked case: residual walks

The residual-walk note is the reference implementation of the rule. Section 9.6 of that note states the boundary and the numeric target ($\kappa_{\text{total}}/\kappa_{\text{floor}} < 0.9$ on a hybrid at fixed B , correct answer on every seed, zero relations failing verification, zero trivial collisions counted as relations). Section 9.7 freezes the baseline table. Section 10.6 classifies every lever against the floor. Sections 11.5–11.6 are an engineering ledger with the class of each step named. Section 11.7 prices the phase the earlier rounds had skipped and corrects the conclusion they implied.³

The relabelling class is not hypothetical. In §10.5 a triple-decomposition oracle cut the walked count from 16.3 to 0.15, a hundredfold move in the number the thread had been tracking, while total cost got 4.6× worse: the count was not reduced, it was paid for per residual instead of per step.

The skipped-phase mistake is the other half. Three rounds priced the relation phase, quoted a crossover with rho at 2^{96} , and left the linear algebra unpriced. When it was measured it grew as $n^{0.68}$ against relations at $n^{1/3}$, so the method’s cost bottoms out near 200× rho around 2^{50} and rises after. The 2^{96} was a statement about one phase and had been read as a statement about the method. That is the §5 mistake of the rule, named so that this paper does not repeat it.

On $E(\mathbb{F}_{p^3})$ at the sizes that fit, rho costs $S \approx 1.3$ and every index-calculus variant built costs between 528× and 4,000× that, with the asymptotically better variant crossing rho only past 2^{230} (extrapolation, exponents in Section 9.3).

³crypto:research/notes/index-calculus/RESEARCH_RESIDUAL_WALKS.md.

2.5 Worked case: the ECC2K-130 pipe

The second worked case is a throughput floor rather than an operation count. Any ECC2K-130 walk that keeps the rho collision structure performs one affine addition per iteration. Affine addition with batched inversion costs a fixed number of field products; a 131×131 carry-less product is six CLMAD. The carry-less unit on the RTX PRO 6000 issues at about 1.6 lane-operations per SM-clock. Those two facts write a floor in billions of iterations per second before any kernel is tuned: 28 B/s needs ≤ 25.4 – 26.9 CLMAD per update including squarings and inversion, and the arithmetic floor is already 28.9 even with squarings moved elsewhere.⁴

Every subsequent GPU row is engineering against that floor. The table walk of Section 6.3 declared 17.0 B/s as its own target before the build; it measured 16.56 and was declined. That is the rule working in the other direction: a pre-declared numeric condition that a later run either meets or does not.

2.6 What does not count

A ratio improved by moving the boundary, changing the unit, or dropping a cost from the budget; a count lowered without the total falling; an extrapolation presented as a measurement; a run whose answer was not checked against the planted secret; wall-clock time as the headline. Extrapolate freely, and mark it as extrapolation with the exponents it rests on. Fit exponents over at least four sizes and report the fit next to the prediction.

3 Pollard rho

Pollard rho is the reference. This section records the walk this library actually runs, the constants it measures, and the GPU and fleet machinery that parallelise it without changing S .

3.1 The walk

The solver is Teske’s r -adding walk [22]. Multipliers $M_i = \alpha_i G + \beta_i H$ are drawn once per job; the state is (Y, a, b) with $Y = aG + bH$; a step is $Y \leftarrow Y + M_{\text{idx}(Y)}$. Distinguished points — hashes with `dp_bits` trailing zero bits — are stored with their (a, b) in a shared table (van Oorschot–Wiener [23]). A second arrival with a different (a, b) gives $(b - b')x = a' - a \pmod{n}$. For composite n the congruence is solved by gcd and every candidate is verified.⁵⁶

r , the number of walks per thread W , and `dp_bits` are chosen automatically from n and the thread count so that neither the multiplier table nor the walk start-ups (each about $3 \log_2 n$ operations) exceed a small fraction of $\sqrt{n}$, and so that the expected number of distinguished points is about 32 per walk, bounded above at 2^{24} table entries. Walks that run $24 \cdot 2^{\text{dp_bits}}$ steps without a distinguished point are restarted.

Threads run batches of W walks; each batch step is one `ca_group_batch_op`, Montgomery’s simultaneous inversion [18]. On one core an affine curve addition costs 232.2 ns with a private inversion and 21.8 ns in batches of 256 (45.88 Mops/s), a $10.6\times$ saving.⁷ A 61-bit $(\mathbb{Z}/p\mathbb{Z})^*$ Montgomery multiply is 5.9 ns (169.4 Mops/s) on the same host.

3.2 Negation map and fruitless cycles

On elliptic curves every point is replaced by the canonical member of $\{Y, -Y\}$ (the one with the smaller Montgomery y), negating (a, b) when needed. The search space shrinks by 2 and

⁴`crypto:ecc2k130/ITERATION-FUNCTION.md`, Åg1.

⁵`src/rho.c`.

⁶`docs/ALGORITHMS.md`, "Pollard rho".

⁷`docs/BENCHMARKS.md`, "Raw group operation throughput".

Table 1: Whole-group discrete logarithm, $S = \text{ops}/\sqrt{n}$, 5 instances per row, 4 threads for `rho-T`. From `ca_bench generic -threads 4`.

grp	bits	BSGS	rho-1	kangaroo	grumpy
$\mathbb{Z}_p^*$	24	1.496	2.066	1.751	1.242
$E(\mathbb{F}_p)$	24	1.496	1.603	1.886	1.211
$\mathbb{Z}_p^*$	32	1.575	1.256	1.412	1.365
$E(\mathbb{F}_p)$	32	1.575	0.791	2.174	1.514
$\mathbb{Z}_p^*$	36	1.441	1.563	1.691	1.351
$E(\mathbb{F}_p)$	36	1.441	1.283	1.768	1.413
$\mathbb{Z}_p^*$	40	1.573	0.887	1.795	1.393
$E(\mathbb{F}_p)$	40	1.573	1.713	1.220	1.251

the expected cost by $\sqrt{2}$. The walk on classes is well defined because `idx` is evaluated on the canonical form.

Fruitless cycles are handled as in Bernstein–Lange–Schwabe [5]:

- *2-cycles*. The look-ahead rule “if $\text{idx}(Y + M_i) = i$ use M_{i+1} instead” (iterated) keeps the walk a deterministic function of Y and reduces the 2-cycle rate from $1/(2r)$ to about $1/(2r)^2$.
- *Longer cycles*. A sliding window compares each new point with a point saved every 64 steps. On a hit the cycle is traversed once, its minimum (by hash) found, and the walk escapes deterministically to $2 \cdot \text{min}$. Because the escape depends only on the cycle, two walks that merged inside a cycle stay merged.

With the negation map, r is allowed to grow to 1024; without it the default is 32.

3.3 Measured S

Table 1 is the whole-group table of `docs/BENCHMARKS.md`: prime-order $(\mathbb{Z}/p\mathbb{Z})^*$ subgroups of safe primes and prime-order curves, five instances per row, S counting every thread’s operations. With five repetitions the rho and kangaroo rows carry about $\pm 25\%$ statistical noise (the standard deviation of one run is about half its mean), so the honest summary is that rho, kangaroo and grumpy giants all sit at $S \sim 1.2\text{--}1.8$, BSGS at 1.5, as theory says.⁸

One-thread rho on 32-bit $\mathbb{Z}_p^*$ is $S = 1.256$; on curves with the negation map the 24–40-bit span is 0.791–1.713. The wall-clock column of the source table (not reproduced) shows the structural facts: table-based methods slow down per operation as the table grows, rho stays at 40–100 Mops/s in $\mathbb{Z}_p^*$ and 10–20 Mops/s on curves, and kangaroo is the fastest per operation because it stores almost nothing.

Interval logarithms of width 2^w inside a $\sim 2^{60}$ ($\mathbb{Z}_p^*$) or $\sim 2^{56}$ (curve) group sit at $S = \text{ops}/\sqrt{w}$ of 1.5–2.8. A separate 400-instance run of grumpy giants gives $S = 1.78$ for intervals and 1.18 for whole groups; the five-instance rows scatter around those values.⁹

3.4 Fitted exponents

S assumes the exponent is $1/2$. The library also measures it. `ca_bench complexity -bits 20,24,28,32,36 -reps 5` fits cost $\sim C \cdot N^\alpha$ in log-log space and reports α next to the theoretical exponent (Table 2).¹⁰

The square-root methods land at $\alpha \sim 0.5$ and the precomputation phases at 0.667 and 0.333. The EC-rho fit $\alpha = 0.428$ is a finite-size artefact (negation-map savings and walk-start constants

⁸`docs/BENCHMARKS.md`, “Whole-group discrete logarithm”.

⁹`docs/ALGORITHMS.md`, “Two grumpy giants and a baby”.

¹⁰`docs/BENCHMARKS.md`, “Measured time complexity”.

Table 2: Fitted exponents, five sizes, five repetitions. The last column is $\text{mean}(\text{cost}/N^{\text{theory}})$, the familiar S , meaningful because the fitted α confirms the exponent.

algorithm	theory	α	R^2	const
BSGS ($\mathbb{Z}_p^*$)	$N^{0.500}$	0.487	0.9993	1.508
rho ($\mathbb{Z}_p^*$)	$N^{0.500}$	0.477	0.9903	1.781
kangaroo ($\mathbb{Z}_p^*$)	$N^{0.500}$	0.526	0.9910	1.843
grumpy ($\mathbb{Z}_p^*$)	$N^{0.500}$	0.502	0.9966	1.264
precomp build ($\mathbb{Z}_p^*$)	$N^{0.667}$	0.647	0.9954	1.192
precomp online ($\mathbb{Z}_p^*$)	$N^{0.333}$	0.347	0.8768	2.781
BSGS (EC)	$N^{0.500}$	0.487	0.9993	1.508
rho (EC)	$N^{0.500}$	0.428	0.9845	1.572
kangaroo (EC)	$N^{0.500}$	0.494	0.9995	1.622
grumpy (EC)	$N^{0.500}$	0.480	0.9956	1.128

Table 3: Host-emulator GPU rho versus CPU rho. S does not grow with the walk count: van Oorschot–Wiener parallelisation is linear in the number of walkers. 3 instances per row.

grp	bits	CPU S	GPU-emu S	walks
$\mathbb{Z}_p^*$	24	1.849	1.895	8
$E(\mathbb{F}_p)$	24	1.707	1.844	8
$\mathbb{Z}_p^*$	28	2.578	1.990	8
$E(\mathbb{F}_p)$	28	2.033	1.504	16
$\mathbb{Z}_p^*$	32	2.045	1.262	40
$E(\mathbb{F}_p)$	32	0.873	1.857	56

on 20–36-bit curves), not a claim that rho beats $N^{1/2}$. Index calculus is absent from the table on purpose: it is $L_p[1/2]$, so no single power-law describes it (Section 9.1).

3.5 GPU kernel, emulator-verified

The same walk runs as a CUDA kernel: tens of thousands of walks, one thread owning eight of them that share a single Fermat inversion (`CA_GPU_W` = 8). A private inversion would be about 99 multiplications per curve step; the batched step is about 18, a $5.5\times$ arithmetic saving.¹¹

The device code is simultaneously valid CUDA C++ and valid C11. A host emulator executes the identical kernel body one thread at a time, so the algorithm is testable without a GPU. Table 3 is that emulator: operation counts are portable, seconds are not a GPU measurement, and no row is a device-throughput result.¹²

Both solvers sit at $S \sim 1.3$ –2.6. A regression — an S that grows with the walk count — would mean the distinguished-point density or the restart policy is wrong, which is what this row watches. Kernel resources from `ptxas` 12.9, no GPU required: 72 registers and 736 B local memory per thread on `sm_70`, 64 registers and the same local memory on `sm_80/89/90`, zero spills on all four.¹³

The host never reads walk state. Distinguished points are written through `atomicAdd`; the host verifies every candidate with a scalar multiplication before returning it.

¹¹`docs/GPU.md`.

¹²`docs/BENCHMARKS.md`, "GPU rho kernel".

¹³`docs/GPU.md`, resource table.

3.6 Fleet protocol

`ca_coord` divides one instance across machines that cannot reach each other. Agents in private subnets dial a coordinator URL; the hub pushes distinguished points and the solution back down the socket the agent opened. What is shared is van Oorschot–Wiener’s observation that parallel rho divides by start point, not by search space: every agent walks the same function (the branch table is derived from the job id) and reports only distinguished points. What is partitioned is the walker index space. Every fact on the wire is verifiable from the job alone: a check-in carries (walker, steps, Y, a, b) with $aG + bH = Y$, which costs the receiver two scalar multiplications to check against the $2^{\text{dp_bits}}$ steps it stands for. The hub is a rendezvous, not an authority.¹⁴

The ECC2K-130 campaign of Section 8 is the same protocol specialised to a Koblitz iteration: 16-bit `-run-id` spaces, 32-byte distinguished-point records, checkpointed iteration totals rather than point counts for rate estimates, and a public dashboard that publishes aggregates only.

4 Square-root neighbours, Cheon, and cross-validation

Rho is the reference. The neighbours are priced in the same unit so that a later claim has somewhere to sit.

4.1 Two grumpy giants and a baby

Bernstein–Lange [3] interleave three walks after shifting so that $x' = x - \ell$ lies in $[0, N)$:

$$B_j = jG, \quad U_i = H' + imG, \quad V_i = 2H' - i(m+1)G.$$

Collisions $B_j = U_i$, $B_j = V_i$, $U_i = V_k$ each give a congruence for x' ; all are taken modulo the group order when it is known, and every candidate is verified. The step $m = \alpha\sqrt{N}$ was chosen by simulating the three walks in exponent space and confirming with the implementation (400 instances): $\alpha = 0.7$ gives $S = 1.18$ on the whole group and $S = 1.78$ on a 2^{22} interval without wrap. The paper’s success probability after $1.5\sqrt{n}$ operations is 0.71875; the simulation gives 0.70 at $n = 4093$. Interleaved BSGS in the same simulation is 1.33, the known $4/3$ figure.¹⁵

4.2 Free precomputation

Bernstein–Lange free precomputation [2, 4] builds, once per group and base, a table of distinguished-point chains and then solves each subsequent target in far fewer operations than a from-scratch square-root search. With $t = \text{round}(\log_2 n/3)$, the table has $\sim n^{1/3}$ chains, the build costs $\sim n^{2/3}$ group operations, and each online target costs $\sim n^{1/3}$. Table 4 is the one-core measurement.¹⁶

The build sits at $P \sim 0.8\text{--}1.6 n^{2/3}$ and the table at $n^{1/3}$ entries, close to the paper’s $1.24 n^{2/3}$ and $n^{1/3}$. The online cost is $T \sim 1\text{--}7 n^{1/3}$ against the paper’s $1.77 n^{1/3}$. The $\sqrt{n}/T$ column is the point of the method: already $7\times\text{--}56\times$ at these sizes, growing as $n^{1/6}$, at the price of a precomputation larger than one $\sqrt{n}$ search and worth it only amortised over many targets in a fixed group. That is a statement about non-uniform security, not a practical attack on a single logarithm.

A fixed-base doubling table, batched inversion, and a 16-byte (fingerprint, exponent) table took the 36-bit curve build from 2.26 s to 0.44 s on one core; four threads then build it in 0.11 s. Class: engineering.

¹⁴`include/cryptanalysis/ca_coord.h`.

¹⁵`docs/ALGORITHMS.md`, grumpy table.

¹⁶`docs/BENCHMARKS.md`, "Discrete logarithm with precomputation".

Table 4: Bernstein–Lange precomputation, one core. P is the one-time build; T the per-target online cost. $\sqrt{n}/T$ is the online speed-up over a from-scratch search.

grp	bits	t	chains	$P/n^{2/3}$	$T/n^{1/3}$	$\sqrt{n}/T$
$\mathbb{Z}_p^*$	24	7	200	1.398	1.926	6.6
$E(\mathbb{F}_p)$	24	8	117	0.892	7.447	1.9
$\mathbb{Z}_p^*$	32	10	787	0.953	2.134	15.0
$E(\mathbb{F}_p)$	32	10	1 558	1.310	2.156	16.7
$\mathbb{Z}_p^*$	36	11	3 039	1.277	0.901	56.4
$E(\mathbb{F}_p)$	36	11	6 145	1.591	2.082	27.4

4.3 GLV folding

On a CM curve the walk is folded by $\langle\psi\rangle$ (order 6 for $j = 0$, 4 for $j = 1728$) against the plain negation-map rho on the same curve. The expected gain is $\sqrt{m/2} - 1.73$ for $j = 0$, 1.41 for $j = 1728$. Measured over 20 instances: `glv-j0-32` lands at $1.71\times$ ($S = 0.988$ against 1.688), `glv-j1728-26` at $1.34\times$. The generic curve is identical either way. Scatter is large because these are small curves; the point is the folding, applied automatically once the j -invariant is recognised.¹⁷

4.4 Cheon

Cheon recovers α from $g, g^\alpha, g^{\alpha^d}$ when d divides $p-1$ [7, 8]. It is not bare ECDLP. With $d \sim \sqrt{p}$ the library measures $1.9\times$ fewer group operations than rho at 32 bits, $4.2\times$ at 40 bits, $18.6\times$ at 48 bits (each Cheon step is an exponentiation, $\sim 1.5 \log_2 p$ operations).¹⁸ The companion records a sharper statement on the $p-1$ and $p+1$ families with fixed-base-comb BSGS: at 40 bits, $S = 0.016$ on $p-1$ ($0.010\times$ matched rho, fitted exponent 0.231 versus rho 0.487); the $p+1$ family fits exponent 0.355 versus rho 0.503. These are advances only against the DLP-with-auxiliary-inputs floor $\Omega(\sqrt{p/d})$, not against the generic ECDLP bound.¹⁹

4.5 Cross-validation against the companion

Two independent implementations in the sibling recover the same constants this library measures. They are cited as walk-checks. The 256-bit CUDA rho and the ECC2K-95 kernel have, per their READMEs, run only through CPU emulation; no device-throughput number is taken from them.

Rho $S = 0.85$ versus BSGS $S = 1.07$. On a 40-bit toy curve ($n = 649\,523\,094\,257$), CPU emulation of the companion kernels, 128 chains, every answer verified as $kG = Q$: rho with the negation map measures $S = 0.85$ (predicted $0.96 \times \sqrt{\pi n/4}$); rho without it $S = 1.42$; BSGS with the negation layout and batched $W = 16$ measures $S = 1.07$ (predicted 1.00), $1.26\times$ the rho reference and $1.07\times$ its own floor. Textbook BSGS without negation is $S = 1.44$ against a predicted 1.41. The batched stepper is engineering (S unchanged, inversions cut by W); the negation layout is the $\sqrt{2}$ the table predicts; nothing here is an advance on rho’s floor.²⁰

Frobenius-class walk, $0.94\times$ and $1.32\times$. The companion Koblitz kernel walks on classes of $\langle\tau, -1\rangle$ with the ECC2K-130 shape $P \leftarrow P + \tau^{j(P)}(P)$, $j = ((g(x)/2) \bmod 8) + 3$, g the normal-basis Hamming weight. On a 21-bit subgroup, 128 walks, 24 trials: mean 251 steps

¹⁷`docs/BENCHMARKS.md`, "GLV endomorphism-accelerated rho".

¹⁸`docs/BENCHMARKS.md`, "Cheon’s attack".

¹⁹`crypto:research/torsion_auxiliary_inputs_20260918/RESULTS.md`.

²⁰`crypto:gpu/ecc/README.md`, "Cost, measured".

against a predicted $\sqrt{\pi r/(4m)} = 268$, i.e. $0.94\times$ the class-walk prediction and $0.20\times$ a negation-only walk. On a 40-bit subgroup, 8 trials: mean 191 840 against 145 129, i.e. $1.32\times$; 59k of the 192k steps are the distinguished-point tail (128 walks $\times$ about 463 steps to a point), and $145k + 59k \approx 204k$. Rho collision times have a standard deviation about equal to their mean, so 24 and 8 trials resolve these to roughly $\pm 20\%$ and $\pm 35\%$. They confirm the $\sqrt{2m}$ speedup; they are too noisy to rank iteration functions.²¹

The same note records that dropping the $/2$ in j collapses the walk to four branches. Measured on the 40-bit curve, 64 walks, 160 trials: 187 831 steps for the four-branch rule against 151 005 for the eight-branch one. That is the Bailey et al. rule, and the one the FPGA core of Section 7.1 implements.

5 ECC2K-130 arithmetic

ECC2K-130 is the Certicom challenge

$$E : y^2 + xy = x^3 + 1 \quad \text{over } \mathbb{F}_{2^{131}},$$

group order $4 \cdot r$ with

$$r = 680564733841876926932320129493409985129$$

a 129-bit prime [1, 6]. Expected work is $2^{60.9}$ iterations of a Frobenius-equivariant walk. This section records the field, the iteration, the golden model, and the small-curve controls. Throughput, FPGA, and the live campaign follow in Sections 6–8.

5.1 Why a type-II optimal normal basis

$2 \cdot 131 + 1 = 263$ is prime and 2 has order 131 modulo 263, so $\mathbb{F}_{2^{131}}$ has a type-II optimal normal basis. With γ a primitive 263rd root of unity the basis is $\beta_i = \gamma^i + \gamma^{-i}$, $i = 1, \dots, 131$, and an element is a 131-bit vector. Three consequences, and every design in both repositories is built on all three:²²

fact	consequence
$\beta_i \beta_j = \beta_{i+j} + \beta_{i-j}$ (indices folded by $\beta_{263-k} = \beta_k$, $\beta_0 = 0$)	a product is a cyclic convolution of two symmetric 263-bit vectors
$\beta_i^2 = \beta_{2i \bmod 263}$	squaring is a permutation — free, it is wiring — so Frobenius costs nothing in this basis
Hamming weight is basis-invariant between this index order and the normal-basis order	the iteration's j and the distinguished-point test read the vector the hardware already holds

None of that is taken on trust. The golden C model (`fpga/model/ecc2k130.c`) builds the field independently and checks the basis, the product rule, the squaring permutation, that 1 is the all-ones vector, and that the trace is the parity of the weight. The curve is identified the same way: $E_0(\mathbb{F}_2)$ has 4 points, so the Frobenius trace is $t = -1$, and the Koblitz recursion $s_m = t s_{m-1} - 2 s_{m-2}$ gives $\#E(\mathbb{F}_{2^{131}}) = 2^{131} + 1 - s_{131} = 4 \cdot r$. The test then checks the two facts that make this *this* curve: $\sigma^2 + \sigma + 2 = 0$ on every point, and $n \cdot P = \mathcal{O}$ for points multiplied by the cofactor.

²¹`crypto:gpu/ecc2k/README.md`, "How this is verified".

²²`fpga/README.md`.

5.2 Three multiplier realisations

A 131-bit product is realised three ways, kept bit-for-bit compatible so a distinguished-point record re-walks under any of them.

1. *Digit-serial cyclic convolution* (this library’s FPGA core). DIGIT bits of the convolution per cycle; inversion is Itoh–Tsujii [15] along 1, 2, 4, 8, 16, 32, 64, 65, 130, eight multiplications, over the core’s one multiplier. Section 7.1.
2. *Bitsliced software* (companion client, historical GPU control). Thirty-two walks in the bits of a word; the original packed-backend audit measured this control at 0.852294 B/s on an RTX PRO 6000.²³
3. *Packed polynomial-basis products* (companion client, and the GPU kernel imported here). Coordinates live in $\mathbb{F}_2[c]/(m)$, $c = \gamma_1$; a 131-bit product is three 64×64 carry-less multiplies — six `clmad` — plus a reduction. Bernstein–Lange conversions to and from the normal basis (where the weight, the selection and the distinguished-point test live) are $O(K \log K)$ shift-and-mask networks. A single-product variant with word-parallel conversions at the boundaries is the first rung of the GPU ledger in Section 6.1.

The companion also has a pipelined Karatsuba multiplier in VHDL for the F2 engine (Section 7.2).

5.3 Two iteration functions

The FPGA core and the historical campaign walk the Bailey et al. [1] map

$$\begin{aligned} j &= ((\text{HW}(x_R) \gg 1) \bmod 8) + 3, & j \in [3, 10], \\ R' &= \sigma^j(R) + R, & \sigma(x, y) = (x^2, y^2), \\ \text{distinguished} &\iff \text{HW}(x) \leq 34 \end{aligned}$$

(campaign cutoff 32 is the live-fleet predicate; Section 8). One step is one affine addition: one inversion, two multiplications and a free Frobenius, since σ^j is wiring. The inversion is eight of about eleven multiplications, so it is 70% of the step — the single most important number in the FPGA design.²⁴

Only x is reported. On a Koblitz curve $-(x, y) = (x, x + y)$, so x is already the negation-class invariant; the Frobenius class is normalised host-side. No coefficients are tracked: a walk is identified by its seed and a collision is resolved by replaying both walks.

The packed GPU client imported into this library walks a different function, an r -adding walk on the classes of $\langle \sigma, -1 \rangle$ (262 points each):

$$\begin{aligned} h(R) &= (\text{HW}(x)/2) \bmod 8, \\ k(R) &= (\sum_e L(e) x_e) \cdot \text{HW}(x)^{-1} \bmod 131, \\ \varepsilon(R) &= \text{bit } p(R) \text{ of } y, \\ R' &= R + (-1)^\varepsilon \sigma^k(T_h), \end{aligned}$$

with $T_h = [a_h]P + [b_h]Q$ and a 131×8 -point table (48.7 KB) in shared memory. k shifts by one under σ and ε flips under negation, so the walk descends to classes with no canonical representative in the loop. Fruitless 2- and 4-cycles are avoided by advancing h against the last four step tags. The two walks share the field, the model and the 32-byte record; they do not share a corpus.²⁵

²³`crypto:ecc2k130/PACKED.md`.

²⁴`fpga/README.md`, "The walk".

²⁵`ecc2k130/README.md`, "The walk".

The exceptional case $\sigma^j(R) = \pm R$ makes the denominator zero. The FPGA core raises `exc` and stops; dividing would produce a point that is not on the curve, which the core would then report as distinguished.

5.4 Golden model and CPU walker

`fpga/model/ecc2k130.c` is the oracle. The FPGA testbenches compare against vectors it emits; the GPU host (`ecc2k130/src/hostcheck.cpp`) re-slices packed words into the same permuted type-II basis and checks equality, not equivalence up to a basis change. `make` in `ecc2k130/` runs 1,904 checks (gcc and clang in CI) covering field arithmetic, Frobenius networks, basis conversions, table selection and the walk, with no CUDA required.²⁶

The companion client’s source compiles for CUDA (32-bit lanes) and for the CPU (64-bit lanes), so the exact code that would run on a GPU is what the CPU test suite exercises. That CPU walker is the verification path for every GPU row below; it is not a throughput claim.

5.5 Small-curve controls, including ECC2K-95

ECC2K-95 is the retired Certicom Koblitz challenge over $\mathbb{F}_{2^{97}}$, group order $4 \cdot r$ with r a 95-bit prime, about 2^{44} iterations, solved by Harley and collaborators in 1998 [14]. $2 \cdot 97 + 1 = 195$ is composite, so there is no type-II optimal normal basis; the companion runs a polynomial-basis backend with the orbit-invariant weight taken through a linear map. The published secret 37837308472231540269443981458 is the check. Harley’s own measurement was $2.16 \cdot 10^{13}$ iterations against a repository prediction of $1.79 \cdot 10^{13}$. The ECC2K-95 CUDA kernel itself has run only through CPU emulation and is not a headline throughput.²⁷²⁸

Toy Koblitz curves with 21-bit and 40-bit subgroups (Section 4.5) close the pipeline to a verified logarithm in seconds. They are the control that the iteration commutes with the class action, not a substitute for the 129-bit group.

6 ECC2K-130 on GPU

The GPU question is a throughput question at a known exponent. The unit is billions of complete scalar iterations per second on one card. The class of every row is engineering: the walk is the same, the collision structure is the same, and the ratio to the generic-group floor does not move. What can move is the ratio to the carry-less-unit floor of Section 2.5.

6.1 The ledger: 0.85 to 14.64 B/s

The original packed-backend audit measured a bitsliced control at 0.852294 B/s on an RTX PRO 6000.²⁹ The historical shipping packed walk, after native `clmad`, compact state and shared Frobenius masks, is 14.637530 B complete iterations/s on the same part (14.106673 B/s while collecting DP34).³⁰ Those two numbers are not a single paired comparison; the rungs between them were measured against adjacent controls, and the early rows of a previous draft of this paper had quoted gains derived from headline numbers rather than from each note’s own pair. Table 5 uses the notes’ stated figures.

Sources, one per rung: single product ³¹; denominator cache ³²; Frobenius networks ³³;

²⁶`ecc2k130/README.md`.

²⁷`crypto:ecc2k130/README.md`.

²⁸`crypto:gpu/ecc2k/README.md`.

²⁹`crypto:ecc2k130/PACKED.md`.

³⁰`crypto:ecc2k130/THROUGHPUT-30B.md`.

³¹`crypto:ecc2k130/SINGLE-PRODUCT.md`.

³²`crypto:ecc2k130/DENOMINATOR-CACHE.md`.

³³`crypto:ecc2k130/FROBENIUS-NETWORK.md`.

Table 5: Packed-backend rungs on one RTX PRO 6000. Each gain is the note’s own paired comparison, not a ratio of non-adjacent headlines. Occupancy changes are named. Class: engineering.

step	from	to	stated gain
bitsliced control	—	0.852294	historical baseline
single product, equal occupancy (2 blocks/SM)	2.956981	3.366572	+13.85%
denominator cache, equal occupancy	3.506509	4.024103	+14.76%
same, retuned to 6 blocks/SM	3.506509	4.094367	+16.76%
Frobenius permutation networks	4.115575	5.811421	+41.21%
polynomial-basis batch chain	5.832342	6.054107	+3.80%
paired products	6.123614	6.215430	+1.50%
<code>clmad</code> + compact state + shared masks	7.110440	14.637530	shipping packed walk

polynomial chain ³⁴; paired products ³⁵; shipping packed walk ³⁶. The polynomial-chain and paired-product comparisons also change occupancy; the notes say so, and so does the table. The compounded remainder between the Frobenius-network confirmation (5.81 B/s) and the shipping packed walk (14.64 B/s) is native carry-less multiplication plus layout work that put the kernel back on the integer pipe, about a $2.5\times$ remaining factor of which the comparable chain of occupancy-controlled rungs above accounts for only a slice. Quoting a single “ $\approx 23\%$ compounded remainder” over the occupancy-matched product/cache/network/chain/pair steps is a narrative of those five notes, not a sixth measurement.

6.2 The floor that rules out 28 B/s

Any walk that keeps the rho collision structure performs one affine addition per iteration. At batch 16 the packed kernel issues 4.81 polynomial products per update; each product is 6 CLMAD; squarings and the inversion chain add more. The arithmetic floor is 28.9 CLMAD/update even with squarings moved to the ALU, and about 1,090–1,180 ALU slots/update. A 28 B/s target on the 188-SM RTX PRO 6000 permits only 25.4–26.9 CLMAD/update and 959–1,013 ALU slots at 100% utilisation. Both budgets are below both floors. 28 B/s on one GPU is therefore not an engineering target; it is below the pipe.³⁷

A later measurement on this library’s imported client, with the carry-less unit issuing at 1.62 lane-operations per SM-clock, prices an iteration at 33.1 `clmad` and writes a floor of 22.3 B/s at 2.42 GHz with the unit never idle. No subsequent row changes that count.³⁸

6.3 Table walk, declined at 16.56 against 17.0

A redesigned r -adding table walk removes the per-step Frobenius and basis-conversion work. On the same host and geometry as the shipping σ -walk it measured 16.560337 B/s against 14.412634 B/s (+14.9%, six alternating repetitions, every paired ratio between 1.146 and 1.169). In the audited 385,024-worker geometry the pair is 16.351936 against 14.975448 (+9.2%). 300 of 300 device reports re-walked; 64 of 64 planted logs were recovered on GF(2^{41}).³⁹

³⁴ `crypto:ecc2k130/POLYNOMIAL-CHAIN.md`.

³⁵ `crypto:ecc2k130/PAIRED-PRODUCTS.md`.

³⁶ `crypto:ecc2k130/THROUGHPUT-30B.md`.

³⁷ `crypto:ecc2k130/ITERATION-FUNCTION.md`, §1.

³⁸ `ecc2k130/README.md`, "Boundary and measurement".

³⁹ `crypto:ecc2k130/benchmarks/table-walk/comparison.json`.

Table 6: Modal GPU survey, automatic workers, packed recipe. Unit: billion complete iterations/s. Class: engineering.

GPU	B/s	/ 6000	it/\$ (T)
RTX PRO 6000, 385k workers	15.116	1.000	17.95
RTX PRO 6000, automatic	14.470	0.957	17.19
B200	8.840	0.585	5.09
L40S	8.696	0.575	16.04
H200	7.632	0.505	6.05
B300	7.575	0.501	3.84
H100	7.540	0.499	6.87
A100-80GB	4.632	0.306	6.67
L4	2.561	0.169	11.54
T4 (no clmad)	0.541	0.036	3.30

The note set 17.0 B/s as its own target in §4.4 *before* the build. Both medians missed it (by 2.6% and 3.8%), so by the rule declared there the change is engineering that did not pay: the shipping walk stays the campaign default and the table walk stays available behind `WALK_TABLE=1`.

This library’s imported client then took that table walk further under a different geometry (`__launch_bounds__(512,1)`, denominators rebuilt from the step tag, inlined products, pipelined forward pass) and measured 20.077–20.080 B/s, 0.90 of its 22.3 B/s floor, with 300/300 device reports re-walked on the golden model.⁴⁰ That configuration is not the live campaign; mixing it with the 14.64 B/s historical ledger without naming the geometry would be the unit mistake.

6.4 Per-GPU survey

The campaign default remains the RTX PRO 6000. Modal also rents T4, L4, A10, L40S, A100, H100, H200, B200 and B300. Table 6 is one packed walk, automatic workers, three valid `-bench` repeats, on every catalog SKU. No row beat 15.115792 B/s. Closest other parts: B200 at 0.585 $\times$, L40S at 0.575 $\times$.⁴¹

The walk is integer/carry-less-bound, not tensor-core-bound. LLM SKU rankings do not transfer. A collecting 6000 at 14.1 B/s is 16.75 T it/\$ at Modal list; L40S is the cheapest other SKU at 16.04 and still loses.

6.5 This library’s imported client

The device code in `ecc2k130/include/` is imported from the companion at the configuration that measured 20.08 B/s. This library’s part is the host: a bridge to the golden model, a re-walk of every device report, and a `make test` that does not need a GPU. The measured run re-walked 300 of 300 reports. CI compiles the kernel and reports registers and spills; no runner has the card, so only a card reproduces 20.08. The live campaign remains the companion’s σ -walk at ~ 14.1 B/s collecting.

7 ECC2K-130 on FPGA

Two FPGA designs, two measurements, one field. This library’s core is a digit-serial walker that simulates with open tools and reports area in yosys LUT4; it has no vendor place-and-route, so

⁴⁰`ecc2k130/README.md`.

⁴¹`crypto:ecc2k130/benchmarks/modal-gpus/SURVEY.md`.

Table 7: Digit-serial ONB core, Icarus Verilog. Area is yosys mapped to 4-input LUTs, not a placed-and-routed device figure. No fmax, so no iterations/s.

DIGIT	cyc/mul	cyc/inv	cyc/step	LUT4 (1 core)
1	132	—	1 342	7 301
2	67	—	692	7 561
4	34	289	362	7 941
8	18	—	202	8 607
16	10	—	122	10 045

it has no iterations-per-second figure. The companion engine is a pipelined Karatsuba walker measured on an AWS F2. They are not the same number in different units.

7.1 Digit-serial core

`fpga/rtl/ecc2k130_step.v` walks $R \leftarrow \sigma^j(R) + R$ with one multiplier per core, Itoh–Tsujii over that multiplier, and back-pressure that stops the walk rather than dropping a point. The collector is round-robin: a starved core stops walking, so a fixed-priority arbiter would quietly reduce a fleet of cores to one. Table 7 is Icarus Verilog, cycles per operation, every row’s RTL having passed the full vector comparison against the golden model.⁴²

Two and four cores at DIGIT = 4 map to 15,705 and 31,437 LUT4: cores scale linearly, they share nothing. From DIGIT 1 to 16 the cycles per step fall $11\times$ while the area grows $1.38\times$, so the area–time product improves about $8\times$, with the caveat that a wider digit lengthens the combinational path and only vendor timing analysis says whether the clock survives it.

What is deliberately not quoted: iterations per second, points per day, “break time”, or a comparison against the GPU. All of those need an fmax this flow cannot produce. The host tool `fpga/host/ec2k` starts, walks, verifies and merges distinguished-point files; `verify` replays every point the board reported against the golden model.

7.2 Companion engine on F2

The companion VHDL engine uses a pipelined Karatsuba multiplier in the same permuted type-II basis, a step unit that shares one inversion across a batch of walks ($5 + 5/W$ multiplies per step), and a sequencer that reports distinguished points. Testbenches are checked against the client’s field model. On an AWS `f2.6xlarge` (VU47P):

image	clocks/step	G steps/s
48 engines	5.31	3.01
64 engines	5.31	4.01
80 engines	5.31	5.02
96 engines, 333 MHz	5.31	6.02
112 engines, 375 MHz	—	8.12
128 engines, 333 MHz	5.17	8.25
120 engines, 333 MHz	—	7.73

The 128-engine image at 333 MHz holds 8,249 M steps/s, i.e. 8.25 G steps/s, with distinguished points sampled from each engine checked against the client’s reference walk and no drops in the listed run.⁴³⁴⁴ Figure 1 draws the measured F2 images against engine count. The engineering estimate before the board was 5–12 B/s per VU47P, central estimate about 8 B/s;

⁴²`fpga/README.md`, “Numbers”.

⁴³`crypto:hdl/ecc2k130/README.md`.

⁴⁴`crypto:hdl/ecc2k130/aws/README.md`.




Figure 1: Companion ECC2K-130 FPGA engine on AWS F2. The 128-engine image at 333 MHz is 8.25 G steps/s. Linear scaling with engine count is the prediction of a design that shares nothing between engines except the host port. This library’s digit-serial core is not on this figure: it has no fmax.

the measurement landed on that central estimate and is a device number, not the estimate reprised.⁴⁵

FPGA workers speak the same client contract as the GPUs and write the same 32-byte records, so a collision between an FPGA trail and a GPU trail is an ordinary merge.

8 ECC2K-130 campaign and structural negatives

The live campaign is a σ -walk, not the 20.08 B/s table-walk corpus. Its configuration is `crypto:ecc2k130/aws/campaign.json`: `dpWeight = 32`, batch 16, 256-thread blocks, two blocks/SM, 385,024 workers = 6,160,384 logical walks, 1,024 steps/launch, `maxIters = 230`. Native carry-less multiplication is what moved the campaign onto CUDA 13.3 and onto batch 16.⁴⁶

8.1 Cutoff 32, 2^{28.41} iterations per point

An earlier fleet-wide ratio of $7.93 \cdot 10^{15}$ checkpointed iterations against 31.6 M uploaded points had been quoted as 2^{27.9} iterations per point. Neither counter was clean: the checkpoint sum multiplied every slot’s per-walk base by one campaign constant, and the point total mixed in the first two slots, which walked at weight 34. A ratio of a mis-weighted sum to a mixed count is not a measurement of the running cutoff.

The calibration that replaced it reads, from each live worker, the client’s own two counters — iterations this run, computed on the grid the client built itself, and distinguished points over the same interval. Frozen at 2026-09-17 18:47 UTC, 132 workers (Table 8):⁴⁷

Expected work remains $2^{60.9} = 2.15 \cdot 10^{18}$ iterations. At weight 32 the expected corpus is about 2^{32.5} records, 0.19 TB; about 40 points/s, 3.4 M points/day, 0.11 GB/day per GPU at the 14.11 B/s DP34 collection rate of the audited packed σ -walk.⁴⁸

⁴⁵`crypto:ecc2k130/FPGA-CEILING.md`.

⁴⁶`crypto:ecc2k130/aws/campaign.json`.

⁴⁷`crypto:ecc2k130/benchmarks/dp-interval/README.md`.

⁴⁸`crypto:ecc2k130/aws/README.md`.

Table 8: Iterations per distinguished point at cutoff 32. One report per $2^{28.41} = 3.56 \cdot 10^8$ iterations, the same on every family, per-worker spread 0.04 in the exponent.

family	workers	iterations	points	it/point
RTX PRO 6000 (g7e)	38	$5.606 \cdot 10^{15}$	15.7 M	$2^{28.41}$
L40S (g6e)	34	$2.723 \cdot 10^{15}$	7.6 M	$2^{28.41}$
RTX PRO 4500 (g7)	48	$3.081 \cdot 10^{15}$	8.6 M	$2^{28.41}$
CPU	12	$2.804 \cdot 10^{14}$	0.79 M	$2^{28.41}$
all	132	$1.169 \cdot 10^{16}$	32.8 M	$2^{28.41}$

8.2 Certification and the 124 M-point snapshot

The storage protocol uses checksummed manifests, pinned worker/replay binaries and lease-fenced checkpoint pointers. Every walk seed is derived from a 16-bit `-run-id`; two contributors who pick different ids never repeat each other’s trail. A distinguished point is a 32-byte record; a collision between any two corpora is solved by reloading them and replaying both walks. The public dashboard publishes aggregates only: not point keys, coefficients, seeds, or database URLs. Checkpointed iteration totals — not just point counts — are the conservative fleet-rate estimator.

The “ ~ 124 M points” figure is a dated storage-recovery snapshot, not the expected collision corpus. After ingesting 63,351,803 rows the store went from 67.6 M points to 124.3 M, with zero failed or outstanding objects and no collision.⁴⁹ No collision is the expected state: 124 M points at cutoff 32 is about $2^{47.8}$ iterations, a fraction of $2^{60.9}$.

8.3 Structural negatives

Several alternatives were measured and did not pay. They are listed because a later thread that rediscovers them should find the measurement rather than the folklore.

- *A 28 B/s one-addition walk on one GPU.* Below the carry-less-unit floor (Section 6.2).
- *The table walk as campaign default.* Measured 16.56 B/s against a pre-declared 17.0; declined (Section 6.3).
- *A Modal SKU that beats the RTX PRO 6000.* Surveyed; none did (Table 6).
- *Bitsliced rewrite as the path to 20 B/s.* The generated bitsliced multiplier’s instruction ceiling does not reach it; the packed path with `clmad` is what did.⁵⁰
- *Index calculus on the Koblitz curve as an ECC2K-130 solver.* The boundary ledger’s best Koblitz row at a 39-bit subgroup is $S = 58.8$, $300\times$ signed-Frobenius rho (Section 9.2). Meet-in-the-middle point decomposition at ECC2K-130-like (m, ℓ) sits 35–51 bits over the per-PDP budget.⁵¹
- *Quoting $2^{27.9}$ as the cutoff-32 interval.* Accounting error; replaced by $2^{28.41}$ (Table 8).

Each is engineering, relabelling, or accounting. None is an advance on the generic-group floor, which for this problem is still $\sqrt{\pi r}/(2 \cdot 262)$ class-walk steps.

⁴⁹`crypto:ecc2k130/aws/README.md`, ingest paragraph.

⁵⁰`crypto:ecc2k130/THROUGHPUT-CEILING.md`.

⁵¹`experiments/pdp-scaling/README.md`.

Table 9: Linear sieve in $(\mathbb{Z}/p\mathbb{Z})^*$, safe primes $p = 2q + 1$, 4 threads, this machine. “rho ops” is $1.3\sqrt{q}$. Rows 60 and 62 are the extension.

bits	fb	unk.	rels	sieve s	LA s	pre s	rho ops
32	53	354	1 290	0.002	0.013	0.015	$4.3 \cdot 10^4$
40	125	926	1 919	0.001	0.043	0.045	$6.8 \cdot 10^5$
48	303	2 504	3 657	0.003	0.170	0.173	$1.1 \cdot 10^7$
56	669	6 270	7 794	0.013	0.793	0.806	$1.7 \cdot 10^8$
60*	1 007	10 008	11 396	0.014	1.614	1.629	$7.0 \cdot 10^8$
62*	1 284	12 785	14 675	0.014	2.697	2.712	$1.4 \cdot 10^9$

* extension, this machine; not in the published table.

9 Index calculus over prime fields

Index calculus is the algorithm whose exponent is supposed to beat one half. In $(\mathbb{Z}/p\mathbb{Z})^*$ it does, near 56 bits on this library. On $E(\mathbb{F}_p)$ and $E(\mathbb{F}_{p^3})$ it does not, at any size that fits. A genus-3/4 Jacobian is the positive control that shows the unit can register a crossover when one exists.

9.1 Linear sieve in $(\mathbb{Z}/p\mathbb{Z})^*$

The library implements the Coppersmith–Odlyzko–Schroeppel linear sieve [9]. With $H = \lceil \sqrt{p} \rceil$ and $J = H^2 - p$,

$$(H + c_1)(H + c_2) = J + (c_1 + c_2)H + c_1 c_2 \pmod{p},$$

and the right-hand side is $O(C\sqrt{p})$ for $|c_i| \leq C$, so it is B -smooth far more often than a random residue. Unknowns are the logs of the primes up to B and of the $2C+1$ numbers $H+c$. Heuristic complexity $L_p[1/2, 1]$. Random-exponent collection ($L_p[1/2, 2]$) is kept as an independent check of the linear algebra.⁵²⁵³

Linear algebra is modulo $p-1 = \prod q^e$. For $q \leq 2^{24}$ the factor-base logs are computed directly with Pohlig–Hellman in the small subgroups. For each larger q the relation matrix is solved modulo q : structured Gaussian elimination, then Lanczos on $A^T D A$ with a random diagonal D (LaMacchia–Odlyzko [17]), Hensel-lifted for $e > 1$. Every factor-base logarithm is verified ($\zeta^{\log} = q \pmod{p}$) before it is used. Individual logarithms are randomised smoothness trials of $h\zeta^e$; at these sizes a few hundred trial divisions suffice, so there is no descent stage.

Table 9 is a re-run of `ca_bench ic -threads 4` on this machine, extended from the published 32–56-bit table through 60 and 62 bits. The 32–56 columns match the published factor-base and unknown counts; wall-clock differs because the host differs. The 60 and 62-bit rows did not appear in the published table and are marked as this extension.⁵⁴⁵⁵

The sieve is essentially free; the cost is Lanczos, which grows as (unknowns)². At the published Z_p^* rate of 169.4 Mops/s, rho at 56 bits is about 1.0 s of multiply-equivalents against 0.81 s of index-calculus precomputation on this host, confirming the published crossover near 56 bits. At 60 and 62 bits the same conversion gives about 4.1 s and 8.3 s of rho against 1.63 s and 2.71 s of sieve-plus-Lanczos: the gap widens, as $L_p[1/2]$ against $N^{1/2}$ says it should. Individual logarithms remain negligible (147–1,310 trial divisions).

This is an $L_p[1/2, 1]$ implementation. The state of the art above roughly 100 bits is the number field sieve ($L_p[1/3]$); NFS is out of scope for a 64-bit toolkit.

⁵²`src/indexcalc.c`.

⁵³`docs/ALGORITHMS.md`, “Index calculus”.

⁵⁴`paper/ic-rerun.txt`.

⁵⁵`docs/BENCHMARKS.md`, Index calculus in $(\mathbb{Z}/p\mathbb{Z})^*$.

Table 10: $E(\mathbb{F}_{p^3})$ residual walks, full pipeline. S includes relations, filtering and linear algebra. Every run recovered the planted logarithm.

p	n	variant	S	S/ρ	LA share
271	$2^{24.2}$	dense	2 696	$2,328\times$	0.05%
271	$2^{24.2}$	Wiedemann	2 830	$2,443\times$	0.19%
271	$2^{24.2}$	+ large primes	4 634	$4,000\times$	0.02%
523	$2^{27.1}$	dense	1 989	$1,680\times$	0.10%
1 039	$2^{30.1}$	dense	1 236	$903\times$	0.31%
2 083	$2^{33.1}$	dense	898	$528\times$	1.22%
2 083	$2^{33.1}$	Wiedemann	974	$574\times$	1.61%
2 083	$2^{33.1}$	+ large primes	3 378	$1,989\times$	0.03%

9.2 Elliptic curves: the negative, phase by phase

For elliptic curves over prime fields no index calculus is known that beats rho at any cryptographic size [10, 13, 20]. The sibling’s boundary ledger puts three regimes — a generic prime-field curve, a random binary curve, and a Koblitz curve — on the same axis, every phase charged, every logarithm verified.⁵⁶

Headline cells, full pipeline:

regime	best variant	S	vs rho	vs floor
prime, 23.4-bit subgroup	MITM $m = 3$	56.8	$14.5\times$	$64.1\times$
random binary, 24.4-bit	MITM $m = 3$	206	$89\times$	$232\times$
Koblitz, 39-bit subgroup	orbit-column MITM	58.8	$300\times$	$425\times$

Matched rho on those instances: $S = 3.93$ (prime), $S = 2.31$ (binary), signed-Frobenius $S = 0.20$ (Koblitz). Fitted full-pipeline exponents: prime MITM $m = 3$ at 0.60; binary MITM $m = 3$ at 0.62; Koblitz signed-orbit MITM at 0.54 overall, or 0.51 on cofactor- ≤ 4 rungs. All are above one half. No cell is below matched rho.

Oracle rounds that price only the decomposition of one residual are stage diagnostics. They do not enter the S column and are not reported as speedups.

9.3 $E(\mathbb{F}_{p^3})$ residual walks

The residual-walk note is the worked negative (Section 2.4). The prime-field route loses because it has no proper additive subspace; the S_3 oracle pays per base element. On $E(\mathbb{F}_{p^3})$ a factor base that is an $\mathbb{F}_p$ -subspace lets a residual be a walk, and the counting floor of §9.6 applies.

A MITM triple oracle at $p = 271$ to 2083 ($n = 2^{24.2}$ to $2^{33.1}$) measures $S = 1,224 \rightarrow 3,805$, i.e. $1,057\times \rightarrow 2,240\times$ rho, fitted total exponent 0.69 against a prediction of $2/3$. Building a base-independent S_4 solve left an initial C_3 of 4.7–4.9 million $\mathbb{F}_p$ -multiplications per residual; the relation-only crossover extrapolated to 2^{111} . Engineering in §§11.5–11.6 reduced C_3 to 879,000, a $5.5\times$ improvement, and the relation-only crossover became 2^{96} . That 2^{96} is a stage diagnostic.

Section 11.7 restores full accounting. Sequential Wiedemann over $\mathbb{Z}/n\mathbb{Z}$, multiplications modulo $n \approx p^3$ charged as 9 $\mathbb{F}_p$ multiplications, every run correct (Table 10).⁵⁷

Fitted exponents over the four sizes: dense total $n^{0.32}$, Wiedemann linear algebra $n^{0.68}$ ($N^{2.00}$ in the number of unknowns), double large primes total $n^{0.44}$ near $4/9$. Relations cost $n^{1/3}$ and the linear algebra $n^{0.68}$ at best, so S falls as $n^{-1/6}$ while the linear-algebra term rises as $n^{+0.18}$. Extrapolating the two measured terms, S bottoms out at ≈ 265 around $n \approx 2^{50}$ — still $\approx 200\times$ rho — and rises after. Double large primes cost $1,989\times$ rho at 33 bits; the

⁵⁶crypto:research/notes/index-calculus/RESEARCH_IC_BOUNDARY_LEDGER.md.

⁵⁷crypto:research/notes/index-calculus/RESEARCH_RESIDUAL_WALKS.md, Å§11.7.

Table 11: Joux–Vitse pair-only versus the S_4 hybrid, full pipeline. At 33 bits the per-residual constant has improved about $580\times$ while residual growth wins: relabelling.

p	n	S	S/rho	residuals
271	$2^{24.2}$	427	$534\times$	1 513
523	$2^{27.1}$	1 331	$951\times$	1 557
1 039	$2^{30.1}$	1 655	$1,273\times$	1 810
2 083	$2^{33.1}$	1 659	$1,037\times$	1 857

Table 12: Hyperelliptic index calculus versus distinguished-point rho. “corr” renormalises rho’s walk to $\sqrt{\pi/2}$. Toy-sized, $p \leq 251$.

g	p	N	S_{IC}	S_ρ	corr. ratio
3	61	124 459	1.51	2.30	0.73
3	101	364 747	1.12	1.90	0.64
3	151	1 180 351	0.79	1.98	0.52
3	211	4 620 611	0.53	1.96	0.38
4	31	239 753	1.77	1.90	0.95
4	41	333 041	1.41	1.86	0.80
4	61	16 790 591	0.28	2.17	0.21

extrapolated whole-method crossover is past 2^{230} . Both crossovers are extrapolations and are marked as such.

9.4 Joux–Vitse pair-only

Decompositions into $k - 1$ points [16] cut the per-residual Groebner cost dramatically and grow the residual count. Table 11 is §11.9 of the residual-walk note.⁵⁸

Residuals fit $n^{0.67}$, total $n^{0.72}$, so S grows as $n^{+0.22}$. Pair-only is the best local cell below roughly $n = 2^{28.5}$; at 33 bits it is relabelling. Rho still costs $S \approx 1.3$ and the best row in the table is 427.

9.5 Genus 3/4 Jacobian, the positive control

Adleman–DeMarrais–Huang / Enge index calculus on a hyperelliptic Jacobian of genus g has predicted cost $O(g^2 q^{2-2/g+\varepsilon})$ against rho’s $O(\sqrt{q^g}) = O(q^{g/2})$. The ratio S_{IC}/S_ρ should therefore fall as $N^{1/g-1/2}$: flat at genus 2, slope $-1/6$ at genus 3, slope $-1/4$ at genus 4. Figure 2 and Table 12 are the sibling’s round-four/five measurements, every row a verified k .⁵⁹

Fitted slopes against $\log_2 N$: genus 2 at -0.027 ($R^2 = 0.24$, no trend, prediction 0); genus 3 at -0.154 ($R^2 = 0.98$, prediction $-1/6$); genus 4 at -0.275 ($R^2 = 0.96$, prediction $-1/4$). The law holds. Genus 3 at $p = 61$ and $p = 101$ already sits at $S_{\text{IC}}/S_\rho = 0.66$ and 0.61 in the raw column; genus 4 at $p = 61$, $N \approx 2^{24}$, reaches 0.21 after the walk renormalisation. That is the crossover this paper’s unit was built to see, and it appears on Jacobians, not on elliptic curves.

10 Related work

Pollard’s rho method [19] and Shanks’ baby-step giant-step [21] are the generic algorithms. Teske’s r -adding walk [22] is the mixing rule this library uses; van Oorschot–Wiener [23] is the parallelisation; Bernstein–Lange–Schwabe [5] is the correct use of the negation map. Two

⁵⁸crypto:research/notes/index-calculus/RESEARCH_RESIDUAL_WALKS.md, §11.9.

⁵⁹crypto:research/notes/index-calculus/RESEARCH_HYPERELLIPTIC_IC_RHO.md.




Figure 2: Hyperelliptic positive control. The ratio to distinguished-point rho falls at genus 3 and genus 4 with the predicted slopes. No measured end-to-end elliptic variant of Sections 9.2 and 9.3 has crossed this line.

grumpy giants and a baby [3] and free precomputation [2, 4] are priced as neighbours, not as replacements. Wiener–Zuccherato and Duursma–Gaudry–Morain [11, 24] folded rho by the automorphism group; GLV [12] is the constructive counterpart that this library applies automatically on $j = 0$ and $j = 1728$. Cheon [7, 8] is a different problem — DLP with auxiliary inputs — and is labelled as such.

The linear sieve [9] and Lanczos over finite fields [17] are the $(\mathbb{Z}/p\mathbb{Z})^*$ solver. Semaev summation polynomials [20], Gaudry’s index calculus for abelian varieties of small dimension [13], Diem’s analysis of the elliptic discrete logarithm [10], and Joux–Vitse pair decompositions [16] are the elliptic side. The contribution here is not a new algorithm; it is a measurement in one unit, with every phase priced, against a floor that cannot be tuned away.

ECC2K-130 is Bailey et al. [1], building on Harley’s ECC2K-95 [14] and Certicom’s challenge list [6]. Itoh–Tsuji inversion [15] and Montgomery’s simultaneous inversion [18] are the two inversion tricks the designs spend most of their area and most of their GPU time on. The FPGA and GPU numbers in Sections 6 and 7 are this programme’s, cited to their artefacts; they are not a claim to have broken the challenge.

The reporting rule — boundary, one unit, four classes, phases priced separately, extrapolations marked — is the sibling study library’s standing convention, written to stop a thread from spending weeks improving a headline constant while the ratio to the floor stays flat. Wesolowski’s supersingular-isogeny result in time $p^{1/3+o(1)}$ is the target-result profile that convention is aimed at: exponent-first, heuristics numbered, phases booked as per-attempt-cost times inverse success probability. This paper is a measurement paper, not that theorem; the profile is named so the negative elliptic result is read as a finished thread, not as a pause.

11 Conclusion

Three statements survive the accounting.

First, Pollard rho on this library is a measured constant, not a slogan: $S \approx 1.2$ – 1.8 across 24–40 bits, with the negation map, Teske’s r -adding walk, Bernstein–Lange–Schwabe cycle handling, and one inversion per 256 additions all inside the count. Independent implementations in the companion recover the same constant (rho $S = 0.85$ against BSGS $S = 1.07$; a Frobenius-class walk at $0.94\times$ and $1.32\times$ of its own prediction). Nothing here is an advance on the generic-group

floor.

Second, ECC2K-130 is an arithmetic problem at a known exponent. The GPU ledger from a bitsliced control at 0.85 B/s to a packed walk at 14.64 B/s is engineering against a carry-less-unit floor that rules out 28 B/s on one GPU. A table walk measured 16.56 B/s against a pre-declared 17.0 and was declined. A digit-serial FPGA core in this library takes 362 cycles/step at 7,941 LUT4; a companion engine on an AWS F2 holds 8.25 G steps/s. The live fleet walks at cutoff 32, one report per $2^{28.41}$ iterations, with about 124 M points and no collision. Structural alternatives (index calculus on the Koblitz curve, a 28 B/s one-addition walk, a SKU that beats the RTX PRO 6000) were measured and did not pay.

Third, index calculus over $(\mathbb{Z}/p\mathbb{Z})^*$ crosses rho near 56 bits and pulls away. Index calculus on $E(\mathbb{F}_p)$ and $E(\mathbb{F}_{p^3})$ does not: at the sizes that fit, every end-to-end elliptic variant costs between $14.5\times$ and $4,000\times$ matched rho, and the asymptotically better $E(\mathbb{F}_{p^3})$ hybrid bottoms near $200\times$ rho around 2^{50} and rises. The genus-3/4 Jacobian is the positive control that shows the unit can register a crossover when one exists.

The constant in front of $\sqrt{n}$ is the right quantity. On the curves that matter it has not yet fallen below one.

A Reproduction commands

Commands below are from the two repository roots. They reproduce the tables, not the GPU or F2 measurements: those need the named card or image. Every answer the library returns is checked against a planted secret inside the tool; a row that is not ok is not a result.

A.1 This library

```
cmake -S . -B build -DCMAKE_BUILD_TYPE=Release
cmake --build build -j
ctest --test-dir build

# Table 1, Table 2 (S and fitted exponents)
./build/ca_bench generic --threads 4
./build/ca_bench complexity --bits 20,24,28,32,36 --reps 5
./build/ca_bench ops
./build/ca_bench interval
./build/ca_bench precomp --threads 1
./build/ca_bench glv
./build/ca_bench cheon
./build/ca_bench gpu          # host emulator; no device required

# Table 5, including the 60/62-bit extension
./build/ca_bench ic --bits 32,40,48,56,60,62 --threads 4

# Golden model and FPGA vectors (no vendor tools)
make -C fpga test          # or: fpga/scripts/run_sim.sh
make -C ecc2k130          # 1,904 checks, no CUDA
```

Sources: docs/BENCHMARKS.md, docs/ALGORITHMS.md, fpga/README.md, ecc2k130/README.md.

A.2 Sibling study library

```
# Rho / BSGS walk-check (CPU emulation of the kernels)
make -C gpu/ecc test
```

```

make -C gpu/ecc2k test

# Boundary ledger (the three-regime elliptic table)
cargo build --release --bin ic
./target/release/ic boundary --oracles --out /tmp/ledger.json

# Hyperelliptic positive control
cargo run --release --example hyperelliptic_ic_vs_rho
python3 scripts/hyperelliptic_exponent_fit.py

# Residual-walk note: the commands are in
# research/notes/index-calculus/RESEARCH_RESIDUAL_WALKS.md §11
# Frozen scoreboard (read, do not recompute):
# docs/index-calculus-scoreboard.html

```

GPU throughput rows need an RTX PRO 6000 (or the named Modal SKU) and CUDA ≥ 13.3 for `clmad`. The F2 row needs the 128-engine image in `hdl/ecc2k130/aws/`. The campaign counters are frozen in `ecc2k130/benchmarks/dp-interval/raw/`.

A.3 This paper

```

cd paper
make          # ecdlp.pdf
make check    # unresolved refs, overfull boxes, page count
make arxiv    # arxiv-submission.tar.gz
make eprint   # eprint-submission.zip

```

Restore authors with `cp AUTHORS.example AUTHORS`, edit, and make `deanonymize`.

References

- [1] Daniel V. Bailey, Lejla Batina, Daniel J. Bernstein, Peter Birkner, Joppe W. Bos, Hsieh-Chung Chen, Chen-Mou Cheng, Gauthier van Damme, Giacomo de Meulenaer, Junfeng Fan, Tim Güneysu, Christof Güraçar, Thorsten Kleinjung, Tanja Lange, Nele Mentens, Ruben Niederhagen, Christof Paar, Francesco Regazzoni, Peter Schwabe, Leif Uhsadel, and Anthony Van Herrewege. Breaking ECC2K-130. *IACR Cryptology ePrint Archive*, (2009/541), 2009.
- [2] Daniel J. Bernstein and Tanja Lange. Computing small discrete logarithms faster. In *INDOCRYPT 2012*, pages 317–338. Springer, 2012.
- [3] Daniel J. Bernstein and Tanja Lange. Two grumpy giants and a baby. In *ANTS X*, volume 1 of *The Open Book Series*, pages 87–111, 2013.
- [4] Daniel J. Bernstein and Tanja Lange. Non-uniform cracks in the concrete: the power of free precomputation. In *ASIACRYPT 2013*, pages 321–340. Springer, 2013.
- [5] Daniel J. Bernstein, Tanja Lange, and Peter Schwabe. On the correct use of the negation map in the Pollard rho method. In *Public Key Cryptography – PKC 2011*, pages 128–146. Springer, 2011.
- [6] Certicom. Certicom ECC challenge. 1997. <https://www.certicom.com/content/certicom/en/the-certicom-ecc-challenge.html>.

- [7] Jung Hee Cheon. Security analysis of the strong Diffie–Hellman problem. In *EUROCRYPT 2006*, pages 1–11. Springer, 2006.
- [8] Jung Hee Cheon. Discrete logarithm problems with auxiliary inputs. *Journal of Cryptology*, 23(3):457–476, 2010.
- [9] Don Coppersmith, Andrew M. Odlyzko, and Richard Schroepel. Discrete logarithms in $\text{GF}(p)$. *Algorithmica*, 1:1–15, 1986.
- [10] Claus Diem. On the discrete logarithm problem in elliptic curves. *Compositio Mathematica*, 147(1):75–104, 2011.
- [11] Iwan Duursma, Pierrick Gaudry, and François Morain. Speeding up the discrete log computation on curves with automorphisms. In *ASIACRYPT 1999*, pages 103–121. Springer, 1999.
- [12] Robert P. Gallant, Robert J. Lambert, and Scott A. Vanstone. Faster point multiplication on elliptic curves with efficient endomorphisms. In *CRYPTO 2001*, pages 190–200. Springer, 2001.
- [13] Pierrick Gaudry. Index calculus for abelian varieties of small dimension and the elliptic curve discrete logarithm problem. *Journal of Symbolic Computation*, 44(12):1690–1702, 2009.
- [14] Robert Harley. Elliptic curve discrete logarithms: ECC2K-95. 1998. Announcement of the solution of Certicom ECC2K-95.
- [15] Toshiya Itoh and Shigeo Tsujii. A fast algorithm for computing multiplicative inverses in $\text{GF}(2^m)$ using normal bases. *Information and Computation*, 78(3):171–177, 1988.
- [16] Antoine Joux and Vanessa Vitse. Elliptic curve discrete logarithm problem over small degree extension fields. In *ASIACRYPT 2013*, 2013. See also Joux–Vitse, “Cover and decomposition index calculus on elliptic curves having a tiny degree of restriction”, 2011–2013.
- [17] Brian A. LaMacchia and Andrew M. Odlyzko. Solving large sparse linear systems over finite fields. In *CRYPTO 1990*, pages 109–133. Springer, 1991.
- [18] Peter L. Montgomery. Speeding the Pollard and elliptic curve methods of factorization. *Mathematics of Computation*, 48(177):243–264, 1987.
- [19] J. M. Pollard. Monte Carlo methods for index computation (mod p). *Mathematics of Computation*, 32(143):918–924, 1978.
- [20] Igor Semaev. Summation polynomials and the discrete logarithm problem on elliptic curves. *IACR Cryptology ePrint Archive*, (2004/031), 2004.
- [21] Daniel Shanks. Class number, a theory of factorization, and genera. In *1969 Number Theory Institute*, volume 20 of *Proceedings of Symposia in Pure Mathematics*, pages 415–440. American Mathematical Society, 1971.
- [22] Edlyn Teske. On random walks for Pollard’s rho method. *Mathematics of Computation*, 70(234):809–825, 2001.
- [23] Paul C. van Oorschot and Michael J. Wiener. Parallel collision search with cryptanalytic applications. *Journal of Cryptology*, 12(1):1–28, 1999.
- [24] Michael J. Wiener and Robert J. Zuccherato. Faster attacks on elliptic curve cryptosystems. In *Selected Areas in Cryptography*, pages 190–200. Springer, 1999.